


MiniGPT-Dense-125M-v1 architecture

Decoder-only GPT: token stream $\rightarrow$ embeddings $\rightarrow$ 12 Transformer blocks $\rightarrow$ tied vocabulary head



Parameters: 123,963,648 • context length: 512 • GPT-2 tokenizer vocabulary: 50,257