


Hybrid Muon optimizer split

Same model parameters, two update rules: matrix-aware Muon for hidden 2D weights, AdamW for fallback parameters



Muon is not replacing AdamW everywhere. The point of the paper implementation is the split.

This makes the benchmark fair: same model, same data, same seed, same token budget — only optimizer logic changes.