


MiniGPT-Dense-125M-v1 architecture and optimizer split



Hybrid Muon applies matrix-aware updates to hidden 2D matrices; AdamW handles embeddings, norms, biases/vectors.